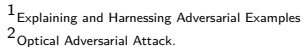


Contribution

We propose Oliva , a novel framework that accelerates **branch-and-bound** neural network verification by incorporating an **order leading exploration** of sub-problems, characterized by:

- A **greedy order** to prioritize sub-problems likely to contain true counterexamples (Oliva^{GR}).
- A **simulated annealing-inspired** approach for **stochastically balancing** exploration and exploitation (Oliva^{SA}).

This framework significantly accelerates the verification process by effectively **navigating** the sub-problem space, and **falsifying** the verification task.



Over-Approximation for Neural Network Verification

Specification: $\Phi \wedge \Psi$ Input: $\Phi := x_1 \in [-1, 1] \wedge x_2 \in [-1, 1]$ Output: $\Psi = (O > 0)$

Network: f





Figure: DeepPoly


- 3 Lowerbound $\hat{p} : \min O = \text{LPSolver}(\Phi \wedge f \wedge \Psi)$
 $\hat{p} = -2.7$ obtained by conservative over-approximation of active functions (i.e., ReLU) via linear solver and cannot verify the specification.

³An Abstract Domain for Certifying Neural Networks





Specification Distance

Specification: $\Phi \wedge \Psi$ Input: $\Phi := x_1 \in [-1, 1] \wedge x_2 \in [-1, 1]$ Output: $\Psi = (O > 0)$

Network: f








$\hat{p} = -2.7$ is the **specification distance** of the overapproximation.



Spurious Counterexample

Specification: $\Phi \wedge \Psi$ Input: $\Phi := x_1 \in [-1, 1] \wedge x_2 \in [-1, 1]$ Output: $\Psi = (O > 0)$

Network: f





Now, it cannot verify the specification.
Branch-and-bound specification tree splitting is needed!


Specification Tree

Specification: $\Phi \wedge \Psi$ Input: $\Phi := x_1 \in [-1, 1] \wedge x_2 \in [-1, 1]$ Output: $\Psi = (O > 0)$

Network: f



Precise splitting ReLU r_1





Over-approximations for ReLUs r_2 r_3 r_4

Specification Tree

Specification: $\Phi \wedge \Psi$ Input: $\Phi := x_1 \in [-1, 1] \wedge x_2 \in [-1, 1]$ Output: $\Psi = (O > 0)$

Network: f





ReLU r_1 is selected to be split. An *atomic proposition* with respect to the selected r_1 neuron is enforced as either $x_1 \geq 0$ (denoted r_i^+) or $x_1 < 0$ (denoted r_i^-).


Specification Tree

Specification: $\Phi \wedge \Psi$ Input: $\Phi := x_1 \in [-1, 1] \wedge x_2 \in [-1, 1]$ Output: $\Psi = (O > 0)$

Network: f




Precise splitting ReLU r_1



Over-approximations for ReLUs r_2 r_3 r_4

ReLU r_1 is selected to be split. An *atomic proposition* with respect to the selected r_1 neuron is enforced as either $x_1 \geq 0$ (denoted r_1^+) or $x_1 < 0$ (denoted r_1^-).





Specification Tree


Specification Tree

Specification: $\Phi \wedge \Psi$ Input: $\Phi := x_1 \in [-1, 1] \wedge x_2 \in [-1, 1]$ Output: $\Psi = (O > 0)$

Network: f




Precise splitting ReLU r_1



Over-approximations for ReLUs r_2 r_3 r_4


ReLU r_1 is selected to be split. An *atomic proposition* with respect to the selected r_1 neuron is enforced as either $x_1 \geq 0$ (denoted r_1^+) or $x_1 < 0$ (denoted r_1^-).



Specification Tree

A specification tree \mathcal{T} consists **nodes** as tuples $\langle \Gamma, \hat{p} \rangle$ that record the subproblems explored. Γ records the split subproblems with the enforced ReLU conditions. Subproblem with the split ReLU (r_1^+) is computed by $\hat{p} = \text{LPSolver}(\Phi \wedge f \wedge \Psi \wedge r_1^+)$, while the rest of the ReLUs are kept approximated.


Existing Approach




Tree grows when continuing to split ReLU r_2 ⁴.

⁴ Branch and Bound for Piecewise Linear Neural Network Verification, J. of Machine Learning Research 2020.

Existing Approach




Tree grows when continuing to split ReLU r_2 ⁴.




Keep branching ReLU r_2 , while ReLUs r_3 and r_4 are kept over-approximation, we find the true counterexample in the end.

⁴ Branch and Bound for Piecewise Linear Neural Network Verification, J. of Machine Learning Research 2020.

Our Insights




Our Insights



From our insights, **counterexamples** are related to two factors:

- The nodes (specification distance \hat{p}) with smaller values are more likely to have a true counterexample.
- The node's depth measures the over-approximation level ($|\Gamma|$): deeper with less approximation

Our Insights




From our insights, **counterexamples** are related to two factors:

- The nodes (specification distance \hat{p}) with smaller values are more likely to have a true counterexample.
- The node's depth measures the over-approximation level ($|\Gamma|$): deeper with less approximation

If we find a true **counterexample** in the subproblem, it is also a **counterexample** for the root problem (step 1).


Limitations of existing approaches and our contribution



- The conventional BaB algorithm ⁴ **exhaustively** verifies the subproblems (each leaf node of BaB tree), which is **inefficient** in discovering counterexamples.
- Our approach explores the **order** of the subproblems to quickly identify true counterexamples (i.e., **eagerly falsify** the verification instance based on the counterexample potential of the subproblem spaces).

⁴Branch and bound for piecewise linear neural network verification. J. of Machine Learning Research 2020.

Contribution I: Inferring the Potentiality of Counterexample




Contribution:

Order-leading verification of subproblems based on counterexample potentiality

$$[\Gamma] = \begin{cases} -\infty & \text{if } \hat{p} > 0 \\ +\infty & \text{true CE} \\ \lambda \frac{[\Gamma]}{K} + (1 - \lambda) \frac{\hat{p}}{\hat{p}_{min}} & \text{otherwise} \end{cases}$$

Contribution I: Inferring the Potentiality of Counterexample




Contribution:

Order-leading verification of subproblems based on counterexample potentiality

$$\llbracket \Gamma \rrbracket = \begin{cases} -\infty & \text{if } \hat{p} > 0 \\ +\infty & \text{true CE} \\ \lambda \frac{\llbracket \Gamma \rrbracket}{K} + (1 - \lambda) \frac{\hat{p}}{\hat{p}_{min}} & \text{otherwise} \end{cases}$$

$\llbracket \Gamma_A \rrbracket = 0.7$
 $\llbracket \Gamma_{B_1} \rrbracket = 0.9$
 $\llbracket \Gamma_{B_2} \rrbracket = 1.1$
 so Γ_{B_2} is selected

Contribution I: Inferring the Potentiality of Counterexample



Contribution:


Order-leading verification of subproblems based on counterexample potentiality

$$\llbracket \Gamma \rrbracket = \begin{cases} -\infty & \text{if } \hat{p} > 0 \\ +\infty & \text{true CE} \\ \lambda \frac{\llbracket \Gamma \rrbracket}{K} + (1 - \lambda) \frac{\hat{p}}{\hat{p}_{min}} & \text{otherwise} \end{cases}$$

$\llbracket \Gamma_A \rrbracket = 0.7$
 $\llbracket \Gamma_{B_1} \rrbracket = 0.9$
 $\llbracket \Gamma_{B_2} \rrbracket = 1.1$
 so Γ_{B_2} is selected

- The nodes ($\llbracket \Gamma_i \rrbracket$) with larger **rewards** (of counterexample potentiality) are more likely to have a true counterexample.


Order Leading Verification Approach (Oliva^{GR})



The conventional BaB algorithm manages "first-come, first-served" processing order, using


- 7 times of bounding, and 2 of branching (only two ReLUs) to find the counterexample.

Order Leading Verification Approach (Oliva^{GR})



The conventional BaB algorithm manages "first-come, first-served" processing order, using



- 7 times of bounding, and 2 of branching (only two ReLUs) to find the counterexample.



Oliva^{GR} considers the **order** of the subproblems (subtrees). Hence, it can speed up the verification process and use

- 5 times of bounding, and 2 times of bounding to find the counterexample.



Contribution II: Simulated Annealing (Oliva^{SA})



Greedily guided by **counterexample potentiality order** may:

- Trapped in **local optima** (similar to “Hill Climbing”).
- unable to achieve **further** improvement.


Contribution II: Simulated Annealing (Oliva^{SA})



Greedily guided by **counterexample potentiality order** may:

- Trapped in **local optima** (similar to “Hill Climbing”).
- unable to achieve **further** improvement.


Contribution II: Simulated Annealing (Oliva^{SA})



Oliva^{SA} with simulated annealing can:

- Possible escape **local optima** and occasionally accept worse solution.
- **Temperature** controlled to achieve **further** improvement by balancing **exploration** and **exploitation**.
- Slowly converge to **exploitation**.


Contribution II: Simulated Annealing (Oliva^{SA})



Oliva^{SA} with simulated annealing can:

- Possible escape **local optima** and occasionally accept worse solution.
- **Temperature** controlled to achieve **further** improvement by balancing **exploration** and **exploitation**.
- Slowly converge to **exploitation**.


Contribution II: Simulated Annealing (Oliva^{SA})



Oliva^{SA} with simulated annealing can:

- Possible escape **local optima** and occasionally accept worse solution.
- **Temperature** controlled to achieve **further** improvement by balancing **exploration** and **exploitation**.
- Slowly converge to **exploitation**.


Oliva^{SA} Child Selection Policy



$$\Delta p \leftarrow \exp \left(\frac{\min R(\Gamma \cdot a) - \max R(\Gamma \cdot a)}{T} \right) \quad \text{s.t. } a \in \{r_k^+, r_k^-\}$$

$$\Gamma^* \leftarrow \Gamma \cdot a^* \quad \text{s.t. } a^* \leftarrow \begin{cases} \text{randomly choose } r_k^+ \text{ or } r_k^- & \text{if } \text{rand}(0, 1) < \Delta p \\ \arg \max_{a \in \{r_k^+, r_k^-\}} R(\Gamma \cdot a) & \text{otherwise} \end{cases}$$


Oliva^{SA} Child Selection Policy



$$\Delta p \leftarrow \exp \left(\frac{\min R(\Gamma \cdot a) - \max R(\Gamma \cdot a)}{T \uparrow} \right) \quad \text{s.t. } a \in \{r_k^+, r_k^-\}$$

$$\Gamma^* \leftarrow \Gamma \cdot a^* \quad \text{s.t. } a^* \leftarrow \begin{cases} \text{randomly choose } r_k^+ \text{ or } r_k^- & \text{if } \text{rand}(0, 1) < \Delta p \downarrow \\ \arg \max_{a \in \{r_k^+, r_k^-\}} R(\Gamma \cdot a) & \text{otherwise} \end{cases}$$

Oliva^{SA} Child Selection Policy




$$\Delta p \leftarrow \exp \left(\frac{\min R(\Gamma \cdot a) - \max R(\Gamma \cdot a)}{T} \right) \quad \text{s.t. } a \in \{r_k^+, r_k^-\}$$

$$\Gamma^* \leftarrow \Gamma \cdot a^* \quad \text{s.t. } a^* \leftarrow \begin{cases} \text{randomly choose } r_k^+ \text{ or } r_k^- & \text{if } \text{rand}(0, 1) < \Delta p \\ \arg \max_{a \in \{r_k^+, r_k^-\}} R(\Gamma \cdot a) & \text{otherwise} \end{cases}$$


Order Leading Verification Approach (Oliva)

Experiment Results with model MNIST-L2 by 241 problem instances



Each point is a verification problem:

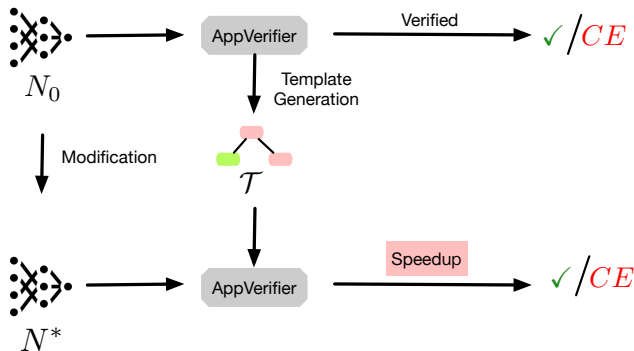
- x-axis: time costs by BaB-baseline
- y-axis: our speedup over BaB-baseline



Distribution of our speedup for violated instances and for certified instances.

Future Work

Order-leading **Incremental** Neural Network Verification - OOPSLA'25





<https://github.com/DeepLearningVerification/Oliva>



Paper

Reference

to be added